

MongoDB Data

Learn about how data is stored in MongoDB.

Introduction

By now, we know that MongoDB is a NoSQL document database and that databases like these have some distinct differences relative to relational databases. In this article, we will dive deeper into the technical aspects of the data that is stored inside of a MongoDB database.

Specifically, we will:

- Dive deeper into how MongoDB stores data via collections and documents
- Gain familiarity with two data file formats: JSON and BSON
- Examine how these data formats are related to data stored in MongoDB

Let's get started!

Collections and Documents

Recall that MongoDB uses the document model. This means that data stored in a MongoDB database resides in a document within a collection. But what does that actually look like? To help better visualize the document model, let's imagine we are using MongoDB to run our camera store. Naturally, we need to keep track of purchases, our customers, etc. Let's break down each layer of the store's database.

At the highest level, we have our database – an instance of MongoDB that contains all the various data our store needs to operate.

Within this instance of MongoDB are collections. Collections are subsets of our data. So, assuming our database contains three types of data – purchase data, inventory, and customer info – each of these would have its own collection.

Within each collection, we store individual records called documents. These documents all belong to that particular subset of our data. So, for example, within the customer collection, we could store personal information about

each of our customers. Every customer would have their own document within the collection.

To summarize, a document is simply a record that stores information about a particular entity. A collection, in turn, is just a group of documents containing similar information. And finally, a MongoDB database is just a number of collections assembled together to store data for a specific use case – in this case running our camera store. This is what the hierarchy would look like visually:

Now that we have a general idea of how data is stored in a MongoDB database, let's take a closer look at documents and how the information they contain is formatted for flexible storage.

Multiple choice

Which of the following correctly describes the structure of a MongoDB database?

A database contains collections; these collections contain documents that store individual records of data.

A database contains documents; within these documents are a number of fields called a collection.

A collection is a group of databases, within these databases are the documents which store our data.

A database contains both collections and documents. A document has a single record of data, whereas a collection has multiple records of data.

Data as JSON

One of the main advantages of using a document database is the flexibility it provides with respect to how data is stored. In the case of MongoDB, this flexibility comes partly from a data format called Javascript Object Notation, or JSON (jay-sahn). [JSON](#) is simply a text format for storing data. Here is what JSON looks like:

```
{  
  "name": "Rodney",  
  "occupation": "photographer",  
}
```

```
"years_of_experience": 7
}
```

In the above example, we have JSON that stores information about a job applicant for a position at our camera store. Note that JSON stores data as what is known as "key-value" pairs, which are always within a pair of curly braces ("{ }"). MongoDB and various online resources also refer to these pairs as "field-value" or "name-value" pairs. For the purposes of this course, we will refer to them as "field-value" pairs. Here is a breakdown of what fields and values mean in the context of JSON:

- **Fields:** A field is a unique identifier for a data point; it tells us what kind of data is being stored. In the job application JSON we just examined, the fields are "name", "occupation", and "years_of_experience". Note that JSON field names must be double-quoted (" "). There are a [number of opinions](#) with respect to how these names should be formatted.
- **Values:** Every field has an associated value. The values are the data points themselves. In our camera store applicant example, the values are "Rodney", "photographer", and 7, each of which corresponds to their respective field. Note that JSON can host a variety of data types in value fields such as strings, numbers, arrays, or even nested data. Additionally, each of these field-value pairs must be separated by a comma (",").

The primary advantage of JSON is readability and flexibility. Data is stored in an easily editable format that is totally comprehensible to humans as well as our computers. However, convenience ultimately comes at a price.

There are three main drawbacks to storing data as JSON:

1. JSON is inefficient from a computational perspective as text is time-consuming to parse.
2. Its readability as text also means that it is not efficient storage-wise. For example, it might be helpful for us to have descriptive names of fields, but they tend to be longer and, for that reason, take up more space.
3. JSON only supports a very limited number of data types – dates, for instance, are not supported natively.

While there are some clear advantages to using JSON to format data, these drawbacks make JSON a poor choice as a primary storage format inside databases. That is why MongoDB invented BSON.

Free response

What are the main advantages and disadvantages of JSON?

Your response

Advantages: readability, flexibility. Drawbacks: time-consuming, not efficient storage-wise, limited number of data types.

Our answer

JSON's usefulness as a format for data comes primarily from the ease with which users can view and modify it. Data is stored in an easily editable format that is totally comprehensible to humans as well as our computers. It can be changed fairly easily with little fuss. However, it has several disadvantages, including slow performance, poor storage efficiency, and limited data types.

Edit Response

BSON - MongoDB's Storage Format

Binary Javascript Object Notation, or BSON (bee-sahn), is the format that MongoDB uses to store data. BSON is different than JSON in three fundamental ways:

1. BSON is not human-readable.
2. BSON is far more efficient storage-wise.
3. BSON supports a number of data formats that JSON does not - like dates.

Our same JSON object from earlier looks like this in BSON:

```
\x00\x00\x00\x02name\x00\xa\x00\x00\x00Rodney\x00\x02occupation\x00\r\x00\x00\x00photographer\x00\x10year_of_experience\x00\xa\x00\x00\x00\x00
```

While it may not be legible, MongoDB wrote the [BSON specification](#) and invented the format to bridge the gap between JSON's flexibility and readability and the required performance of a large database. MongoDB stores data as BSON internally but allows users to create and manipulate database data as JSON. This allows for both efficient data storage and a great developer experience!

Free response

How is BSON different from JSON?

Your response

BSON is not human readable, is more efficient storage-wise, and supports a wider range of data types (for example dates).

Our answer

BSON is unlike JSON in three main ways:

1. BSON does not store data as text and so it is not readable by humans
2. Partly because it doesn't store data as text, it is far more efficient than JSON

3. BSON supports a wider array of data formats natively, like dates for example

Edit Response

Wrap Up

In this article, we learned some of the basics of how data is stored in MongoDB. Specifically, we learned:

- A MongoDB instance can contain many databases, and within a database are collections of similar data. Collections contain individual records called documents that are stored as field-value pairs.
- JSON is a human-readable but inefficient format for storing data; conversely, BSON is not human-readable but is highly efficient.
- MongoDB users can easily store and manipulate data as JSON – even though internally, that data is stored as BSON.

MongoDB's use of the document model, combined with its dual use of JSON and BSON, makes it a highly efficient and usable option for storing data. To continue learning and get a more in-depth overview of MongoDB data, check out the following two guides on JSON/BSON from MongoDB:

- [JSON and BSON Guide](#)
- [Explaining BSON with Examples](#)